

L6a: Data-Driven Minimum-Variance Portfolios

CHEME 5660 Cornell University

Disclaimer and Risks

Training only. This content is for training and informational purposes only. The teaching team does not make, give, or endorse any offer or solicitation to buy or sell securities or derivative products, or provide investment or trading advice or strategy.

Risk. Trading involves risk and is generally not recommended for those with limited resources, experience, or a low risk tolerance. Past performance does not guarantee future success.

Responsibility. You are responsible for your investment decisions based on your financial circumstances, objectives, risk tolerance, and liquidity needs.

Objectives for Today

By the end of this lecture, you should be able to:

- **Formulate portfolio reward and risk:** derive matrix expressions for expected growth and variance, and explain diversification through covariance.
- **Construct minimum-variance portfolios:** compute GMV weights allowing shorts, identify the efficient branch, and explain how long-only constraints, exact growth targets, and growth floors affect the solution.
- **Evaluate estimated allocations:** explain how estimated inputs affect the weights, and evaluate wealth and NPV using prices reserved for testing.

Lectures and Examples

Lecture: [Data-Driven Minimum-Variance Portfolios](#)

Example 1: [Portfolio Weights and Dirichlet Sampling](#)

Draw long-only allocations, compare their estimated growth and risk, and follow selected buy-and-hold portfolios.

Example 2: [Minimum-Variance Portfolios from Data](#)

Estimate inputs from 2014–2024 data and trace the frontier. Compare long-only GMV, equal weights, and an index fund using 2025 prices.

Example 3: [Portfolio Wealth with Multiple Asset GBM](#)

Simulate correlated prices and buy-and-hold wealth; estimate the probability of exceeding a target scaled NPV.

Concept Review: Portfolio Weights

How much of our money should we put into each asset? For M assets, w_i is the fraction of initial wealth invested in asset i :

$$\underbrace{\mathbf{w} = [w_1, \dots, w_M]^\top}_{\text{mix of assets}}, \quad \underbrace{w_i \geq 0}_{\text{long only}}, \quad \underbrace{\sum_{i=1}^M w_i = 1}_{\text{fully invested}}.$$

These constraints define the **allocation simplex**:

- Two assets give a line segment; three assets give a triangle.
- Each corner invests the entire budget in one asset. Equal weights lie at the center.
- With initial wealth $W_0 > 0$, invest $w_i W_0$ dollars in asset i . Holding the purchased shares fixed allows the weights to change with prices.

We first sample feasible allocations, then develop an optimization problem to choose the weights.

Concept Review: Dirichlet Sampling

The [Dirichlet distribution](#) generates fully invested, long-only allocations. Draw $\mathbf{w} \sim \text{Dirichlet}(\alpha)$, with dimensionless $\alpha_i > 0$ and $\alpha_0 = \sum_i \alpha_i$. Each weight has:

$$\boxed{\mathbb{E}[w_i] = \frac{\alpha_i}{\alpha_0}}, \quad \boxed{\text{Var}(w_i) = \frac{\alpha_i(\alpha_0 - \alpha_i)}{\alpha_0^2(\alpha_0 + 1)}}.$$

- The ratios α_i/α_0 set the expected weights; larger α_0 at fixed ratios reduces their spread across draws.
- With equal $\alpha_i = \alpha$: $\alpha < 1$ favors the boundary, $\alpha = 1$ samples uniformly, and large α clusters draws near equal weights.

Finding the minimum variance over all feasible weights requires optimization.

Example 1: [Portfolio Weights and Dirichlet Sampling](#)

Company Profile: AllianceBernstein

AllianceBernstein manages investments for institutional and individual clients.

- **Investment management:** equity, fixed-income, and multi-asset strategies, supported by quantitative research and data science.
- **Managing risk:** the [AB International Low Volatility Equity ETF](#) combines quantitative and fundamental research, emphasizing risk mitigation.
- **Artificial intelligence:** [Andrew Chin](#), Chief Artificial Intelligence Officer and Cornell alumnus, leads the firm's AI strategy. I know Andrew through the [AI in Finance short course](#).

Today's connection: how should we choose portfolio weights to balance expected growth and risk?

Explore: [Student opportunities](#) [Cornell Keynotes conversation with Andrew Chin](#)

Modern Portfolio Theory: Portfolio Reward

For M risky assets, w_i is the initial fraction invested in asset i , with $\sum_i w_i = 1$. Hold the weights fixed when calculating reward and risk. Let $\mathbf{g} = [g_1, \dots, g_M]^\top$ contain the asset growth rates over one observation interval. Their weighted average is:

$$g_p = \sum_{i=1}^M w_i g_i = \mathbf{w}^\top \mathbf{g}.$$

With fixed weights and $\mathbb{E}[g_i] = \mu_{g,i}$, **reward** is:

$$\mathbb{E}[g_p] = \sum_{i=1}^M w_i \mathbb{E}[g_i] = \mathbf{w}^\top \boldsymbol{\mu}_g.$$

The sample-mean vector $\mathbf{g}'$ estimates $\boldsymbol{\mu}_g$. Estimated reward is $\mathbf{w}^\top \mathbf{g}'$ (year⁻¹).

Expected growth is linear in the portfolio weights.

Weighted Growth and Buy-and-Hold Wealth

With fixed share counts, one-period wealth and its log growth rate are given by:

$$\frac{W_{\Delta t}}{W_0} = \sum_{i=1}^M w_i e^{g_i \Delta t},$$

$$\frac{1}{\Delta t} \ln \left(\frac{W_{\Delta t}}{W_0} \right) = \frac{1}{\Delta t} \ln \left(\sum_{i=1}^M w_i e^{g_i \Delta t} \right), \quad W_{\Delta t} > 0.$$

The logarithm acts on a sum, so this generally differs from $g_p = \sum_i w_i g_i$.

Use weighted growth rates to choose weights, then use buy-and-hold wealth to evaluate the investment through time.

Next, we measure how the weighted growth rate varies across outcomes.

Modern Portfolio Theory: Portfolio Risk

We measure **risk** by $\text{Var}(g_p)$. With fixed weights and covariance matrix $\Sigma_g = \text{Cov}(\mathbf{g})$:

$$\begin{aligned}\text{Var}(g_p) &= \text{Cov}\left(\sum_i w_i g_i, \sum_j w_j g_j\right) \\ &= \sum_i \sum_j w_i w_j \text{Cov}(g_i, g_j) = \mathbf{w}^\top \Sigma_g \mathbf{w}.\end{aligned}$$

Individual variances (diagonal entries) and pairwise covariances (off-diagonal entries) both contribute to risk. From data, use $\hat{\Sigma}_g$.

Variance has units year^{-2} . Its square root, $\sigma_{g,p} = \sqrt{\text{Var}(g_p)}$, is the growth-rate standard deviation (year^{-1}).

To see how covariance can reduce risk, consider a two-asset portfolio.

Two-Asset Diversification

Invest fractions w and $1 - w$, with $0 < w < 1$, in two assets with positive standard deviations $\sigma_{g,1}$ and $\sigma_{g,2}$ and correlation ρ_{12} :

$$\text{Var}(g_p) = w^2\sigma_{g,1}^2 + (1 - w)^2\sigma_{g,2}^2 + 2w(1 - w) \underbrace{\rho_{12}\sigma_{g,1}\sigma_{g,2}}_{\text{Cov}(g_1, g_2)}.$$

Correlation determines what mixing achieves:

- At $\rho_{12} = 1$, the standard deviation is the weighted average:
 $\sigma_{g,p} = w\sigma_{g,1} + (1 - w)\sigma_{g,2}.$
- For $\rho_{12} < 1$, it falls below that weighted average.
- If ρ_{12} is below the smaller standard deviation divided by the larger, some allocation has less variance than either asset. For 2 and 4 per year, the cutoff is $\rho_{12} < 0.5$.

Positive correlation can still permit diversification.

Growth Rates and Log Returns

For an observation interval $\Delta t > 0$, asset log returns are $\mathbf{r} = \Delta t \mathbf{g}$. Their weighted average is $r_p = \mathbf{w}^\top \mathbf{r} = \Delta t g_p$:

$$\mathbb{E}[r_p] = \Delta t \mathbb{E}[g_p],$$

$$\text{Var}(r_p) = \Delta t^2 \text{Var}(g_p).$$

Log returns are dimensionless, with covariance $\Sigma_r = \Delta t^2 \Sigma_g$.

Scaling preserves the optimization problem:

- Every portfolio variance is multiplied by the same positive factor Δt^2 , preserving the minimizing weights.
- Scale the target as $r_\star = \Delta t g_\star$ so that the same portfolios meet the target.

Growth rates and log returns therefore give the same optimal weights.

Standard Deviation and Volatility

Growth-rate standard deviation depends on the observation interval. To calculate volatility, use the covariance rate $\mathbf{C} = \Delta t \Sigma_g$ from L5b:

$$\sigma_{g,p} = \sqrt{\mathbf{w}^\top \Sigma_g \mathbf{w}} \quad (\text{year}^{-1}),$$

$$\text{GBM volatility} = \sqrt{\mathbf{w}^\top \mathbf{C} \mathbf{w}} = \sqrt{\Delta t} \sigma_{g,p} \quad (\text{year}^{-1/2}).$$

For daily observations, $\Delta t = 1/252$ years:

$$\hat{\mathbf{C}} = \frac{\hat{\Sigma}_g}{252}.$$

The **risk axis** in this lecture shows $\sigma_{g,p}$, measured in year^{-1} .

With reward and risk defined, we can now choose the weights that minimize variance.

The Global Minimum-Variance Portfolio

We seek the fully invested portfolio with the smallest growth-rate variance. **Short positions are allowed** ($w_i < 0$), and there is **no expected-growth constraint**. Assume Σ_g is symmetric and positive definite.

With $\mathbf{1} \in \mathbb{R}^M$ the vector of ones, the problem is:

$$\begin{aligned} & \underset{\mathbf{w} \in \mathbb{R}^M}{\text{minimize}} && \frac{1}{2} \mathbf{w}^\top \Sigma_g \mathbf{w} \\ & \text{subject to} && \mathbf{1}^\top \mathbf{w} = 1. \end{aligned}$$

The solution and its variance are:

$$\mathbf{w}_{\text{GMV}} = \frac{\Sigma_g^{-1} \mathbf{1}}{\mathbf{1}^\top \Sigma_g^{-1} \mathbf{1}}, \quad \sigma_{g,\text{GMV}}^2 = \frac{1}{\mathbf{1}^\top \Sigma_g^{-1} \mathbf{1}}.$$

The weights depend only on covariance. We now derive this result.

Deriving the GMV Weights

Introduce a multiplier λ for the budget constraint. The factor $1/2$ simplifies the derivative without changing the minimizing weights:

$$\mathcal{L}(\mathbf{w}, \lambda) = \frac{1}{2} \mathbf{w}^\top \Sigma_g \mathbf{w} - \lambda (\mathbf{1}^\top \mathbf{w} - 1).$$

- **Stationarity:** $\Sigma_g \mathbf{w} - \lambda \mathbf{1} = \mathbf{0}$, so $\mathbf{w} = \lambda \Sigma_g^{-1} \mathbf{1}$.
- **Budget:** $1 = \lambda \mathbf{1}^\top \Sigma_g^{-1} \mathbf{1}$, giving $\lambda = (\mathbf{1}^\top \Sigma_g^{-1} \mathbf{1})^{-1}$.
- **Variance:** use $\Sigma_g \mathbf{w} = \lambda \mathbf{1}$ and $\mathbf{w}^\top \mathbf{1} = 1$:

$$\sigma_{g,\text{GMV}}^2 = \mathbf{w}^\top \Sigma_g \mathbf{w} = \lambda \mathbf{w}^\top \mathbf{1} = \lambda = \frac{1}{\mathbf{1}^\top \Sigma_g^{-1} \mathbf{1}}.$$

Positive definiteness makes the objective strictly convex, so these weights give the unique minimum.

The Long-Only Problem

Long-only weights are nonnegative and sum to one. We also require estimated expected growth to reach a **target floor** $g_\star$ (year^{-1}).

Using sample means $\mathbf{g}'$ and estimated covariance $\hat{\Sigma}_g$, solve:

$$\begin{aligned} & \underset{\mathbf{w} \in \mathbb{R}^M}{\text{minimize}} && \mathbf{w}^\top \hat{\Sigma}_g \mathbf{w} \\ & \text{subject to} && (\mathbf{g}')^\top \mathbf{w} \geq g_\star, \\ & && \mathbf{1}^\top \mathbf{w} = 1, \\ & && 0 \leq w_i \leq 1, \quad i = 1, \dots, M. \end{aligned}$$

This is a **convex quadratic program**: the covariance matrix is positive semidefinite and the constraints are linear.

How do these constraints change the solution?

Adding the Portfolio Constraints

Use the same estimated inputs and add the restrictions in order:

1. **First, exclude short positions.** If the original GMV weights are all nonnegative, keep them. Otherwise, solve numerically with $w_i \geq 0$ and $\mathbf{1}^\top \mathbf{w} = 1$. The result is the **long-only GMV portfolio**.
2. **Then, impose the growth floor.** If the long-only GMV portfolio meets g_* , keep its weights. For a higher feasible target, solve again; the weights change and minimum variance increases.

At either step, restricting the portfolio choices cannot lower the minimum variance.

The Target-Growth Frontier

Allow short positions and require $\mathbf{1}^\top \mathbf{w} = 1$. For each **equality target** $\mu_g^\top \mathbf{w} = g_\star$, find the weights with the smallest variance.

Assume Σ_g is positive definite and the asset means are not all equal. Define:

$$\begin{aligned} a &= \mathbf{1}^\top \Sigma_g^{-1} \mathbf{1}, & b &= \mathbf{1}^\top \Sigma_g^{-1} \mu_g, \\ c &= \mu_g^\top \Sigma_g^{-1} \mu_g, & d &= ac - b^2 > 0. \end{aligned}$$


The minimum variance at the target is:

$$\sigma_g^2(g_\star) = \frac{ag_\star^2 - 2bg_\star + c}{d} = \frac{1}{a} + \frac{a}{d} \left(g_\star - \frac{b}{a} \right)^2.$$

Varying $g_\star$ traces the **minimum-variance frontier**. At $g_\star = b/a$, we recover the GMV variance $1/a$; higher or lower targets require more variance.

The Efficient Branch

A portfolio is efficient if no feasible alternative offers at least as much expected growth and no more variance, with one strict improvement.



- p_2 is the GMV portfolio.
- p_1 and p_3 have the same risk; p_1 has higher growth.
- The upper branch, including p_2 , is efficient.

Feasible portfolios lie on or to the right of the frontier; moving right adds variance.

Tracing the Long-Only Frontier

Long-only portfolio growth is a weighted average of the sample means:

$$\min_i g'_i \leq (\mathbf{g}')^\top \mathbf{w} \leq \max_i g'_i.$$

This gives a practical target sweep:

- Start with $g_\star \leq \min_i g'_i$. Every long-only portfolio meets the floor, so the solution is the long-only GMV portfolio.
- Increase the target to trace the efficient branch. Targets above $\max_i g'_i$ are infeasible.
- With the same inputs, long-only constraints cannot reduce variance. The frontiers meet where the shorts-allowed weights are nonnegative.

Let's estimate the inputs, sweep the target, and inspect how the weights change.

Example 2: [Minimum-Variance Portfolios from Data](#)

Estimating Portfolio Inputs

The GMV calculation uses the estimated covariance $\hat{\Sigma}_g$. A target-growth calculation also uses the estimated mean growth rates $\mathbf{g}'$. The weights therefore depend on the data and estimation method.

- For AMD (2014–2024), mean growth estimates are 0.3120 per year from the sample mean and 0.4397 per year from L4b's regression: a difference of 12.8 percentage points per year.
- Underestimating an asset combination's variance can make the optimizer favor it.

Our data example evaluates the selected allocations on 2025 prices, alongside equal weights and an index fund.

Example 3: Portfolio Wealth with Multiple Asset GBM. Hold the fitted inputs and allocation fixed; simulate buy-and-hold wealth and estimate the probability of exceeding a target scaled NPV.

Portfolio Wealth and Scaled NPV

Initial weights $w_i \geq 0$, with $\sum_i w_i = 1$, divide wealth $W_0 > 0$ among the assets. Buy $n_i = w_i W_0 / S_i(0)$ shares and hold them fixed:

$$W_t = W_0 \sum_{i=1}^M w_i \frac{S_i(t)}{S_i(0)},$$
$$\rho_T = \frac{W_T}{W_0} e^{-g_y T} - 1.$$

Here, g_y is the constant benchmark growth rate and T is the holding period in years. We allow fractional shares and omit dividends and trading costs.

Simulated wealth outcomes estimate $\mathbb{P}(\rho_T > \rho_\star)$ for a target scaled NPV $\rho_\star$.

The optimized weights set the initial holdings; wealth and NPV measure how the resulting investment performs through time.

Optional Advanced Material

These examples develop the frontier geometry and input sensitivity further.

Advanced 1: Frontier Geometry and the Two-Fund Theorem

How can two risky frontier portfolios generate the rest of the curve? Derive the weights and vary the mix of two fixed funds. Compare the frontier with a 50% limit on each position's magnitude and with long-only weights.

Advanced 2: How Estimates Affect Portfolio Weights

How much do weights change when we estimate the inputs again? Resample the 2014–2024 observations, vary means and covariance separately, and compare GMV and tangent allocations. Evaluate the resulting long-only portfolios on the same 2025 prices.

Summary

- **Portfolio reward and risk:** we expressed expected growth and variance using the weights, mean growth rates, and covariance.
- **Minimum-variance portfolios:** we derived the GMV weights and used growth targets and weight constraints to trace the efficient frontier.
- **Estimated inputs and performance:** we connected the allocation to its estimates and evaluated buy-and-hold wealth on a separate year of prices.

In L6b, we use a market factor to estimate portfolio inputs and add a risk-free asset.

Next: model asset growth using a shared market factor and firm-specific risk.