


$$a = \arg \max \{a_1, \dots, a_K\}$$


$$a_i \sim \text{Beta}(\alpha_i, \beta_i)$$



Bandit problems

$$\pi(s) = \arg \max_a Q(s, a)$$

$$Q(s, a)$$



Temporal Difference:
Q-learning/SARSA